

Machine Learning 1

Zaida Rodriguez (PID:A59010549)


10/22/2021

#Clustering methods

Kmeans clustering in R is done with the `kmeans()` function Here we make up some data to test and learn with.

```
tmp <- c(rnorm(30, 3), rnorm(30,-3))
data <- cbind(x=tmp, y=rev(tmp))

#the goal of this is to make a data set that has -3 and +3 values in x and y
# x:(-3, +3) and y:(+3, -3)
plot(data)
```



Run `kmeans()` set `k(centers)` to 2 and `nstart` to 20. The thing with Kmeans is you have to tell it how many clusters you want.

```

km <- kmeans(data, centers=2, nstart=2)
km

## K-means clustering with 2 clusters of sizes 30, 30
##
## Cluster means:
##           x           y
## 1 -3.195614  2.958954
## 2  2.958954 -3.195614
##
## Clustering vector:
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
##
## Within cluster sum of squares by cluster:
## [1] 56.88963 56.88963
## (between_SS / total_SS =  90.9 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
## [6] "betweenss"    "size"         "iter"         "ifault"

```

Clustering vector is telling you for which cluster your element belongs to.

Q1. How many points are in each cluster?

```
km$size
```

```
## [1] 30 30
```

Q2. What 'component' of your result object details cluster assignment/membership?

```
km$cluster
```

```
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1

```

Q3. What 'component' of your result object details cluster center?


```
km$centers
```

```
##           x           y
## 1 -3.195614  2.958954
## 2  2.958954 -3.195614

```

Q4. Plot x colored by the kmeans cluster assignment and add clusters as blue points (or by clusters)

```
plot(data, col=km$cluster)
points(km$centers, col="blue", pch=15, cex=2)
```



#Hierarchical clustering

We will use the `hclust()` function on the same data as before and see how this method works.

```
hc <- hclust(dist(data))
hc
```


```
##
## Call:
## hclust(d = dist(data))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

#this function requires you to tell it the distance from each point

`hclust` has a `plot` method

```
plot(hc)
abline(h=7, col="blue")
```

Cluster Dendrogram



```
dist(data)
hclust (*, "complete")
```

```
# the bottom (leaves) are your row names
# it puts what is closest to each other together
# 2 main groups; you can "cut it"
```

To find our membership vector we need to “cut” the tree (dendrogram) and for this we use the `cutree()` function and tell it the height to cut at.


```
cutree(hc, h=7)
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2  
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

We can also use t we want ...

```
grps <- cutree(hc, k=2)
```

```
plot(data, col=grps)
```



Principal Component Analysis (PCA)

PCA OF UK Food data

read the file ine

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
## [1] 17  5
```

```
# there are 17 rows and 5 columns
```

view the first 6 rows

```
head(x)
```

```
##           X England Wales Scotland N.Ireland
## 1      Cheese      105   103      103      66
## 2 Carcass_meat     245   227      242     267
## 3   Other_meat     685   803      750     586
## 4        Fish     147   160      122      93
## 5 Fats_and_oils    193   235      184     209
## 6        Sugars    156   175      147     139
```

fix the row names:

```
rownames(x) <- x[,1]
x <- x[,-1]
x
```

```
##           England Wales Scotland N.Ireland
## Cheese          105   103      103      66
## Carcass_meat    245   227      242     267
## Other_meat      685   803      750     586
## Fish           147   160      122      93
## Fats_and_oils   193   235      184     209
## Sugars          156   175      147     139
## Fresh_potatoes  720   874      566    1033
## Fresh_Veg       253   265      171     143
## Other_Veg       488   570      418     355
## Processed_potatoes 198  203      220     187
## Processed_Veg    360  365      337     334
## Fresh_fruit     1102 1137      957     674
## Cereals          1472 1582     1462    1494
## Beverages        57    73       53      47
## Soft_drinks     1374 1256     1572    1506
## Alcoholic_drinks 375   475      458     135
## Confectionery    54    64       62      41
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

Don’t do things this way, you will overwrite your object. Instead, read it in with row names already(seen below). This is the preferred method as it will not mess with your data and is more robust.


```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
x
```

```
##           England Wales Scotland N.Ireland
## Cheese          105   103      103      66
## Carcass_meat    245   227      242     267
## Other_meat      685   803      750     586
## Fish           147   160      122      93
## Fats_and_oils   193   235      184     209
## Sugars          156   175      147     139
## Fresh_potatoes  720   874      566    1033
## Fresh_Veg       253   265      171     143
```

## Other_Veg	488	570	418	355
## Processed_potatoes	198	203	220	187
## Processed_Veg	360	365	337	334
## Fresh_fruit	1102	1137	957	674
## Cereals	1472	1582	1462	1494
## Beverages	57	73	53	47
## Soft_drinks	1374	1256	1572	1506
## Alcoholic_drinks	375	475	458	135
## Confectionery	54	64	62	41

now create a barplot. By adding “col=rainbow” you can change the color of the bars

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




barplot need to be run as a matrix, since we didn't have a matrix, you can force it right into the fu

Q5. Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

This pairwise plot allows you to compare the 4 countries against each other. The first row compares England (x) and Wales (y), then Scotland (y), then N. Ireland(y) The 2nd row compares Wales (x) and England(y), then N. Ireland (y); and so on and so on. The points that are outliers are visible and those indicate differences in the consumption of the particular food categories depending on the country. The points that lie on the diagonal for a given plot indicate that the values are the same for both countries that are being compared.

```
mycols <- rainbow(nrow(x))
pairs(x, col=mycols, pch=16)
```



> Q6. What is the main differences in between N. Ireland and the other countries of the UK interms of this data-set?

N, Ireland has a greater variation of the food consumption in comparison to England, Wales, and Scotland.

PCA to the rescue!

Not easy to interpret, takes time. Instead try PCA!


Here we will use the base R function for PCA, which is called `prcomp()`. This function wants you to first transpose your data (ie: flip the columns with rows)

```
pca <- prcomp(t(x))
summary(pca)
```

```
## Importance of components:
##
## Standard deviation      324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance  0.6744  0.2905  0.03503 0.000e+00
## Cumulative Proportion  0.6744  0.9650  1.00000 1.000e+00
```



```
plot(pca)
```



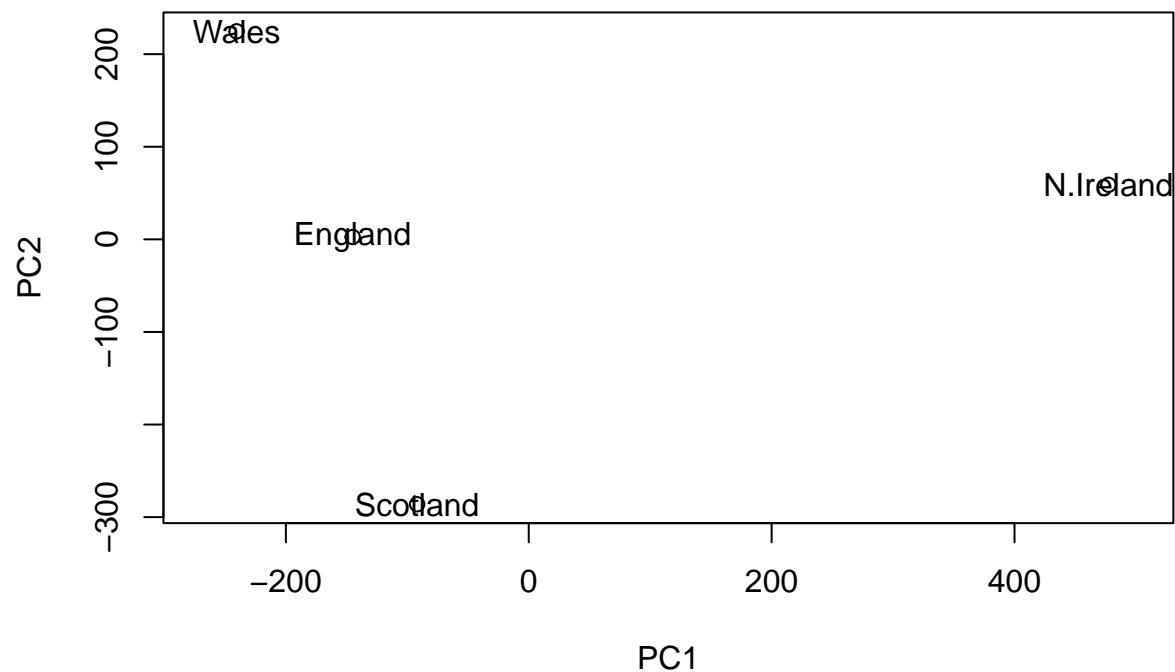
We want a score plot (aka: PCA plot). Basically plot of PC1 vs PC2

```
attributes(pca)
```

```
## $names
## [1] "sdev"      "rotation" "center"    "scale"     "x"
##
## $class
## [1] "prcomp"
```


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
#plot PC1 vs PC2
plot(pca$x[,1:2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1:2], labels=colnames(x))
```




Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at the start of this document.

```
plot(pca$x[,1:2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1:2], labels=colnames(x), col=rainbow(nrow(x)))
```




We can also examine the PCA “loading”, which tells us how much of the original variables contribute to each new PC.

```
barplot(pca$rotation[,1], las=2)
```



> Q9. Generate a similar 'loadings plot' for PC2. What 2 food groups feature prominently and what does PC2 mainly tell us about?

```
barplot(pca$rotation[,2], las=2)
```



the 2 main food groups are fresh potatoes and soft drinks. It tells us how much of the original varia

one more PCA

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##      wt1 wt2 wt3 wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1 439 458 408 429 420 90 88 86 90 93
## gene2 219 200 204 210 187 427 423 434 433 426
## gene3 1006 989 1030 1017 973 252 237 238 226 210
## gene4 783 792 829 856 760 849 856 835 885 894
## gene5 181 249 204 244 225 277 305 272 270 279
## gene6 460 502 491 491 493 612 594 577 618 638
```

Q10. How many genes are there in this data?

```
nrow(rna.data)
```

```
## [1] 100
```

```
# there are 100 genes
```

```
ncol(rna.data)
```

```
## [1] 10
```

```
colnames(rna.data)
```


```
## [1] "wt1" "wt2" "wt3" "wt4" "wt5" "ko1" "ko2" "ko3" "ko4" "ko5"
```

```
pca.rna = prcomp(t(rna.data), scale=TRUE)  
summary(pca.rna)
```


```
## Importance of components:
```

```
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7  
## Standard deviation  9.6237 1.5198 1.05787 1.05203 0.88062 0.82545 0.80111  
## Proportion of Variance 0.9262 0.0231 0.01119 0.01107 0.00775 0.00681 0.00642  
## Cumulative Proportion 0.9262 0.9493 0.96045 0.97152 0.97928 0.98609 0.99251  
##              PC8      PC9      PC10  
## Standard deviation  0.62065 0.60342 3.348e-15  
## Proportion of Variance 0.00385 0.00364 0.000e+00  
## Cumulative Proportion 0.99636 1.00000 1.000e+00
```

```
plot(pca.rna)
```



```
plot(pca.rna$x[,1:2])
text(pca.rna$x[,1:2], labels=colnames(rna.data))
```



PC1 is telling us which genes are changing the most between KO to KO and WT to WT. PC2