

Class_15

Zaida Rodriguez (PID:A59010549)

11/17/2021

Load the contData and colData

it requires 2 things: 1. count data 2. colData (the metadata that tells us about the design of the experiment)

```
library(BiocManager)
```

read in the data

```
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

look at the files

```
head(counts)
```

```
##           SRR1039508 SRR1039509 SRR1039512 SRR1039513 SRR1039516
## ENSG00000000003      723        486        904        445        1170
## ENSG00000000005         0         0         0         0         0
## ENSG00000000419      467        523        616        371        582
## ENSG00000000457      347        258        364        237        318
## ENSG00000000460       96         81         73         66        118
## ENSG00000000938         0         0         1         0         2
##           SRR1039517 SRR1039520 SRR1039521
## ENSG00000000003      1097        806        604
## ENSG00000000005         0         0         0
## ENSG00000000419      781        417        509
## ENSG00000000457      447        330        324
## ENSG00000000460       94        102         74
## ENSG00000000938         0         0         0
```

```
head(metadata)
```

```
##           id      dex celltype      geo_id
## 1 SRR1039508 control   N61311 GSM1275862
## 2 SRR1039509 treated   N61311 GSM1275863
## 3 SRR1039512 control   N052611 GSM1275866
## 4 SRR1039513 treated   N052611 GSM1275867
## 5 SRR1039516 control   N080611 GSM1275870
## 6 SRR1039517 treated   N080611 GSM1275871
```

Side-note: Lets check the correspondence of the metadata and count data setup.

```
metadata$id
```

```
## [1] "SRR1039508" "SRR1039509" "SRR1039512" "SRR1039513" "SRR1039516"  
## [6] "SRR1039517" "SRR1039520" "SRR1039521"
```

or we can use:

```
colnames(counts)
```

```
## [1] "SRR1039508" "SRR1039509" "SRR1039512" "SRR1039513" "SRR1039516"  
## [6] "SRR1039517" "SRR1039520" "SRR1039521"
```

We can use the == things to see if they are the same.

```
metadata$id==colnames(counts)
```

```
## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

This function will look at what you are telling it to and tell if you it is true or not

```
all( c(T,T,F))
```

```
## [1] FALSE
```

For example:

```
all(metadata$id==colnames(counts))
```

```
## [1] TRUE
```

Q1.How many genes ?

```
nrow(counts)
```

```
## [1] 38694
```

Q2.How many control groups are there?

```
head(metadata)
```

```
##           id      dex celltype    geo_id  
## 1 SRR1039508 control   N61311 GSM1275862  
## 2 SRR1039509 treated   N61311 GSM1275863  
## 3 SRR1039512 control   N052611 GSM1275866  
## 4 SRR1039513 treated   N052611 GSM1275867  
## 5 SRR1039516 control   N080611 GSM1275870  
## 6 SRR1039517 treated   N080611 GSM1275871
```

There are 38694 rows/genes in this dataset. There are 3 control gorups in this dataset.

Compare control to treated

First, extract the data from the control groups(columns) in our counts data

```
control.inds <- metadata$dex=="control"  
metadata[control.inds, ]
```

```
##           id      dex celltype      geo_id  
## 1 SRR1039508 control   N61311 GSM1275862  
## 3 SRR1039512 control   N052611 GSM1275866  
## 5 SRR1039516 control   N080611 GSM1275870  
## 7 SRR1039520 control   N061011 GSM1275874
```

Just to get the ID use:

```
control.inds <- metadata[control.inds, ]$id
```

Now we use these ids to access just the control columns for our counts data

```
head(counts[, control.inds])
```

```
##           SRR1039508 SRR1039512 SRR1039516 SRR1039520  
## ENSG000000000003      723        904        1170        806  
## ENSG000000000005         0         0         0         0  
## ENSG000000000419      467        616        582        417  
## ENSG000000000457      347        364        318        330  
## ENSG000000000460       96         73        118        102  
## ENSG000000000938         0         1         2         0
```

this is the count data just for the control groups.

now get the mean of the genes (rows):

```
control.mean <- rowMeans(counts[, control.inds])  
head(control.mean)
```

```
## ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460  
##           900.75           0.00          520.50          339.75           97.25  
## ENSG000000000938  
##           0.75
```

Now do the same for the treated groups:

```
treated.id <- metadata[ metadata$dex=="treated", ] $id  
head(treated.id)
```

```
## [1] "SRR1039509" "SRR1039513" "SRR1039517" "SRR1039521"
```

```
treated.mean <- rowMeans(counts[, treated.id])  
head(treated.mean)
```

```
## ENSG00000000003 ENSG00000000005 ENSG00000000419 ENSG00000000457 ENSG00000000460
##           658.00           0.00           546.00           316.50           78.75
## ENSG00000000938
##           0.00
```

Now combine these two means

```
meancounts <- data.frame(control.mean, treated.mean)
# meancounts
colSums(meancounts)
```


```
## control.mean treated.mean
##    23005324    22196524
```

The colSums shows you the summary for both columns

Now visualize by plotting

```
library(ggplot2)
```

```
plot(meancounts[,1],meancounts[,2], xlab="Control", ylab="Treated")
```




This would benefit from a log transformation.

```
plot(meancounts, log="xy")
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted  
## from logarithmic plot
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted  
## from logarithmic plot
```



We often use log transformations as they make life much nicer in this world.

For example:

```
log2(20/20)
```

```
## [1] 0
```

```
# this indicates that there is no change
```

```
log2(40/20)
```

```
## [1] 1
```

```
log2(10/20)
```

```
## [1] -1
```

```
log2(80/20)
```

```
## [1] 2
```

```
#this is doubling. Gene expression would be doubling
```

```
meancounts$log2fc <- log2(meancounts[, "treated.mean"]/meancounts[, "control.mean"])  
head(meancounts)
```

```
##               control.mean treated.mean      log2fc  
## ENSG000000000003      900.75      658.00 -0.45303916  
## ENSG000000000005         0.00         0.00         NaN  
## ENSG000000000419      520.50      546.00  0.06900279  
## ENSG000000000457      339.75      316.50 -0.10226805  
## ENSG000000000460       97.25       78.75 -0.30441833  
## ENSG000000000938        0.75         0.00        -Inf
```

Now remove the genes that have 0 values

```
head(meancounts[,1:2])
```

```
##               control.mean treated.mean  
## ENSG000000000003      900.75      658.00  
## ENSG000000000005         0.00         0.00  
## ENSG000000000419      520.50      546.00  
## ENSG000000000457      339.75      316.50  
## ENSG000000000460       97.25       78.75  
## ENSG000000000938        0.75         0.00
```

Lets look at which ones have zero values

```
head(meancounts[,1:2] ==0)
```

```
##               control.mean treated.mean  
## ENSG000000000003      FALSE      FALSE  
## ENSG000000000005       TRUE       TRUE  
## ENSG000000000419      FALSE      FALSE  
## ENSG000000000457      FALSE      FALSE  
## ENSG000000000460      FALSE      FALSE  
## ENSG000000000938      FALSE       TRUE
```

Now remove those that have zero values

The which function tells us teh indices of the TRUE entries in a logical vector.

```
which(c(T,F,T))
```

```
## [1] 1 3
```

However, it isn't that useful in the default mode on our type of multi column input ...

```
ind <- which(meancounts[,1:2] == 0, arr.ind=T)[,"row"]
head(ind)
```

```
## ENSG000000000005 ENSG00000004848 ENSG00000004948 ENSG00000005001 ENSG00000006059
##                2                65                70                73                121
## ENSG00000006071
##                123
```

I only care about the rows here (if there is a zero in any column I will exclude this row eventually)

```
to.rm <- unique(sort(ind))
mycounts <- meancounts[-to.rm,]
head(mycounts)
```

```
##                control.mean treated.mean      log2fc
## ENSG00000000003          900.75         658.00 -0.45303916
## ENSG00000000419          520.50         546.00  0.06900279
## ENSG00000000457          339.75         316.50 -0.10226805
## ENSG00000000460           97.25          78.75 -0.30441833
## ENSG00000000971         5219.00        6687.50  0.35769358
## ENSG00000001036        2327.00        1785.75 -0.38194109
```

We now have 21817 genes remaining

```
nrow(mycounts)
```

```
## [1] 21817
```

How many of these genes are upregulated at the log2 fold-change threshold of +2 or greater?

```
sum(mycounts$log2fc > +2)
```

```
## [1] 250
```

What percentage is this?

```
round(sum(mycounts$log2fc > +2) / nrow(mycounts) * 100, 2)
```

```
## [1] 1.15
```

How many of these genes are down regulated at the log2 fold-change threshold of -2 or greater?

```
sum(mycounts$log2fc < -2)
```

```
## [1] 367
```

```
round(sum(mycounts$log2fc < -2) / nrow(mycounts) * 100, 2)
```

```
## [1] 1.68
```

DESeq Analysis

First set it up

```
library(DESeq2)
```

```
## Loading required package: S4Vectors
```

```
## Loading required package: stats4
```

```
## Loading required package: BiocGenerics
```

```
##
```

```
## Attaching package: 'BiocGenerics'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      IQR, mad, sd, var, xtabs
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      anyDuplicated, append, as.data.frame, basename, cbind, colnames,  
##      dirname, do.call, duplicated, eval, evalq, Filter, Find, get, grep,  
##      grepl, intersect, is.unsorted, lapply, Map, mapply, match, mget,  
##      order, paste, pmax, pmax.int, pmin, pmin.int, Position, rank,  
##      rbind, Reduce, rownames, sapply, setdiff, sort, table, tapply,  
##      union, unique, unsplit, which.max, which.min
```

```
##
```

```
## Attaching package: 'S4Vectors'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      expand.grid, I, unname
```

```
## Loading required package: IRanges
```

```
## Loading required package: GenomicRanges
```

```
## Loading required package: GenomeInfoDb
```



```

## Loading required package: SummarizedExperiment

## Loading required package: MatrixGenerics

## Loading required package: matrixStats

##
## Attaching package: 'MatrixGenerics'

## The following objects are masked from 'package:matrixStats':
##
##   colAlls, colAnyNAs, colAnys, colAveragesPerRowSet, colCollapse,
##   colCounts, colCummaxs, colCummins, colCumprods, colCumsums,
##   colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,
##   colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,
##   colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,
##   colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,
##   colWeightedMeans, colWeightedMedians, colWeightedSds,
##   colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAveragesPerColSet,
##   rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,
##   rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,
##   rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,
##   rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,
##   rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,
##   rowWeightedMads, rowWeightedMeans, rowWeightedMedians,
##   rowWeightedSds, rowWeightedVars

## Loading required package: Biobase

## Welcome to Bioconductor
##
##   Vignettes contain introductory material; view with
##   'browseVignettes()'. To cite Bioconductor, see
##   'citation("Biobase)"', and for packages 'citation("pkgname)"'.

##
## Attaching package: 'Biobase'

## The following object is masked from 'package:MatrixGenerics':
##
##   rowMedians

## The following objects are masked from 'package:matrixStats':
##
##   anyMissing, rowMedians

dds <- DESeqDataSetFromMatrix(countData=counts,
                              colData=metadata,
                              design=~dex)

## converting counts to integer mode

```

```
## Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
## design formula are characters, converting to factors
```

```
dds
```

```
## class: DESeqDataSet
## dim: 38694 8
## metadata(1): version
## assays(1): counts
## rownames(38694): ENSG000000000003 ENSG000000000005 ... ENSG00000283120
## ENSG00000283123
## rowData names(0):
## colnames(8): SRR1039508 SRR1039509 ... SRR1039520 SRR1039521
## colData names(4): id dex celltype geo_id
```

Run the DESeq analysis pipeline.

```
dds <- DESeq(dds)
```

```
## estimating size factors
```

```
## estimating dispersions
```

```
## gene-wise dispersion estimates
```

```
## mean-dispersion relationship
```

```
## final dispersion estimates
```

```
## fitting model and testing
```

```
res <- results(dds)
res
```


```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 38694 rows and 6 columns
##      baseMean log2FoldChange    lfcSE      stat      pvalue
##      <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003  747.1942    -0.3507030  0.168246 -2.084470  0.0371175
## ENSG000000000005    0.0000         NA         NA         NA         NA
## ENSG000000000419  520.1342     0.2061078  0.101059  2.039475  0.0414026
## ENSG000000000457  322.6648     0.0245269  0.145145  0.168982  0.8658106
## ENSG000000000460   87.6826    -0.1471420  0.257007 -0.572521  0.5669691
## ...           ...           ...           ...           ...
## ENSG00000283115   0.000000         NA         NA         NA         NA
## ENSG00000283116   0.000000         NA         NA         NA         NA
## ENSG00000283119   0.000000         NA         NA         NA         NA
## ENSG00000283120   0.974916    -0.668258    1.69456 -0.394354  0.693319
## ENSG00000283123   0.000000         NA         NA         NA         NA
##                padj
```

```
##                               <numeric>
## ENSG000000000003  0.163035
## ENSG000000000005      NA
## ENSG000000000419  0.176032
## ENSG000000000457  0.961694
## ENSG000000000460  0.815849
## ...
## ENSG00000283115      NA
## ENSG00000283116      NA
## ENSG00000283119      NA
## ENSG00000283120      NA
## ENSG00000283123      NA
```

A Volcano Plot


This is a very common

```
plot(res$log2FoldChange, res$padj)
```



Run it but take the log

```
plot(res$log2FoldChange, -log(res$padj))
# add a line
abline(v=c(-2,2), col="yellow")
abline(h=-log(0.05), col="blue")
```



Add extra information by annotating gene names and colors. Adding meaningful gene names to our dataset allows us to make sense of what is going on here.

```
library("AnnotationDbi")
```

```
## Warning: package 'AnnotationDbi' was built under R version 4.1.2
```

```
library("org.Hs.eg.db")
```

```
##
```

```
columns(org.Hs.eg.db)
```

```
## [1] "ACCNUM"      "ALIAS"       "ENSEMBL"     "ENSEMBLPROT" "ENSEMBLTRANS"
## [6] "ENTREZID"    "ENZYME"      "EVIDENCE"     "EVIDENCEALL"  "GENENAME"
## [11] "GENETYPE"    "GO"          "GOALL"        "IPI"          "MAP"
## [16] "OMIM"        "ONTOLOGY"    "ONTOLOGYALL" "PATH"         "PFAM"
## [21] "PMID"        "PROSITE"     "REFSEQ"       "SYMBOL"       "UCSCKG"
## [26] "UNIPROT"
```

Here we map to “SYMBOL” the common gene name that the world understands and wants.

```
res$symbol <- mapIds(org.Hs.eg.db,
                    keys=row.names(res),
                    keytype="ENSEMBL",
                    column="SYMBOL",
                    multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
head(res$symbol)
```

```
## ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
##      "TSPAN6"      "TNMD"      "DPM1"      "SCYL3"      "C1orf112"
## ENSG000000000938
##      "FGR"
```

```
head(res)
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 7 columns
##      baseMean log2FoldChange    lfcSE      stat    pvalue
##      <numeric>    <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003 747.194195  -0.3507030  0.168246 -2.084470 0.0371175
## ENSG000000000005  0.000000      NA      NA      NA      NA
## ENSG000000000419 520.134160   0.2061078  0.101059  2.039475 0.0414026
## ENSG000000000457 322.664844   0.0245269  0.145145  0.168982 0.8658106
## ENSG000000000460  87.682625  -0.1471420  0.257007 -0.572521 0.5669691
## ENSG000000000938  0.319167  -1.7322890  3.493601 -0.495846 0.6200029
##      padj      symbol
##      <numeric> <character>
## ENSG000000000003 0.163035      TSPAN6
## ENSG000000000005      NA      TNMD
## ENSG000000000419 0.176032      DPM1
## ENSG000000000457 0.961694      SCYL3
## ENSG000000000460 0.815849      C1orf112
## ENSG000000000938      NA      FGR
```

Lets save our results to date

```
write.csv(res, file="allmyresults.csv")
```

We will be merging some files using the `merge()` function

PATHWAY ANALYSIS

Let's try to bring some biology insight back into our first analysis

```
library(pathview)
```

```
## #####
## Pathview is an open source software package distributed under GNU General
## Public License version 3 (GPLv3). Details of GPLv3 is available at
## http://www.gnu.org/licenses/gpl-3.0.html. Particullary, users are required to
## formally cite the original Pathview paper (not just mention it) in publications
## or products. For details, do citation("pathview") within R.
##
## The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG
## license agreement (details at http://www.kegg.jp/kegg/legal.html).
## #####
```

```
library(gage)
```

```
##
```

```
library(gageData)
```

Visualize the first rows of data

```
data(kegg.sets.hs)
head(kegg.sets.hs, 2)
```

```
## $'hsa00232 Caffeine metabolism'
## [1] "10" "1544" "1548" "1549" "1553" "7498" "9"
##
## $'hsa00983 Drug metabolism - other enzymes'
## [1] "10" "1066" "10720" "10941" "151531" "1548" "1549" "1551"
## [9] "1553" "1576" "1577" "1806" "1807" "1890" "221223" "2990"
## [17] "3251" "3614" "3615" "3704" "51733" "54490" "54575" "54576"
## [25] "54577" "54578" "54579" "54600" "54657" "54658" "54659" "54963"
## [33] "574537" "64816" "7083" "7084" "7172" "7363" "7364" "7365"
## [41] "7366" "7367" "7371" "7372" "7378" "7498" "79799" "83549"
## [49] "8824" "8833" "9" "978"
```

Before we can use KEGG we need to get our gene identifiers in the correct format for KEGG, which is ENTREZ format in this case.

```
head(res)
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 7 columns
##      baseMean log2FoldChange      lfcSE      stat      pvalue
##      <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003 747.194195 -0.3507030 0.168246 -2.084470 0.0371175
## ENSG000000000005  0.000000          NA          NA          NA          NA
## ENSG000000000419 520.134160  0.2061078 0.101059  2.039475 0.0414026
## ENSG000000000457 322.664844  0.0245269 0.145145  0.168982 0.8658106
```

```
## ENSG00000000460 87.682625 -0.1471420 0.257007 -0.572521 0.5669691
## ENSG00000000938 0.319167 -1.7322890 3.493601 -0.495846 0.6200029
##          padj          symbol
##          <numeric> <character>
## ENSG00000000003 0.163035      TSPAN6
## ENSG00000000005      NA      TNMD
## ENSG00000000419 0.176032      DPM1
## ENSG00000000457 0.961694      SCYL3
## ENSG00000000460 0.815849      C1orf112
## ENSG00000000938      NA      FGR
```

```
head(rownames(res))
```

```
## [1] "ENSG00000000003" "ENSG00000000005" "ENSG00000000419" "ENSG00000000457"
## [5] "ENSG00000000460" "ENSG00000000938"
```

```
res$entrez <- mapIds(org.Hs.eg.db,
                     keys=row.names(res),
                     keytype="ENSEMBL",
                     column="ENTREZID",
                     multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
res$genename <- mapIds(org.Hs.eg.db,
                       keys=row.names(res),
                       keytype="ENSEMBL",
                       column="GENENAME",
                       multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

The main **gage()** function requires a named vector of fold changes, where the names of the values are the Entrez gene IDs.

Note that we used the `mapIds()` function above to obtain Entrez gene IDs (stored in `res$entrez`) and we have the fold change results.

```
foldchanges <- res$log2FoldChange
head(foldchanges)
```

```
## [1] -0.35070302      NA 0.20610777 0.02452695 -0.14714205 -1.73228897
```

Assign names to this vector that are the gene IDs that KEGG wants

```
names(foldchanges) = res$entrez
head(foldchanges)
```

```
##          7105          64102          8813          57147          55732          2268
## -0.35070302      NA 0.20610777 0.02452695 -0.14714205 -1.73228897
```

#not meaningful so you could change to symbol, but not necessary at this time because...

Now we are ready for the `gage()` function. Get the results

```
keggres = gage(foldchanges, gsets=kegg.sets.hs)
View(keggres)
```

We can look at the `attributes()` of this or any R object.

```
attributes(keggres)
```

```
## $names
## [1] "greater" "less"    "stats"
```

```
head(keggres$less)
```

```
##                                     p.geomean stat.mean
## hsa05332 Graft-versus-host disease    0.0004250461 -3.473346
## hsa04940 Type I diabetes mellitus     0.0017820293 -3.002352
## hsa05310 Asthma                       0.0020045888 -3.009050
## hsa04672 Intestinal immune network for IgA production 0.0060434515 -2.560547
## hsa05330 Allograft rejection          0.0073678825 -2.501419
## hsa04340 Hedgehog signaling pathway   0.0133239547 -2.248547
##                                     p.val      q.val
## hsa05332 Graft-versus-host disease    0.0004250461 0.09053483
## hsa04940 Type I diabetes mellitus     0.0017820293 0.14232581
## hsa05310 Asthma                       0.0020045888 0.14232581
## hsa04672 Intestinal immune network for IgA production 0.0060434515 0.31387180
## hsa05330 Allograft rejection          0.0073678825 0.31387180
## hsa04340 Hedgehog signaling pathway   0.0133239547 0.47300039
##                                     set.size      exp1
## hsa05332 Graft-versus-host disease    40 0.0004250461
## hsa04940 Type I diabetes mellitus     42 0.0017820293
## hsa05310 Asthma                       29 0.0020045888
## hsa04672 Intestinal immune network for IgA production 47 0.0060434515
## hsa05330 Allograft rejection          36 0.0073678825
## hsa04340 Hedgehog signaling pathway   56 0.0133239547
```

The `pathview()` function will add our genes to a KEGG pathway

```
pathview(gene.data=foldchanges, pathway.id="hsa05310")
```

```
## 'select()' returned 1:1 mapping between keys and columns
```

```
## Info: Working in directory /Users/zaidarodriguez/Desktop/UCSD/Fall2021/BGGN213/bggn213_github/Class1
```

```
## Info: Writing image file hsa05310.pathview.png
```